

Stage 3: Prompt Library & Specifications

Stage Three: Prompt Library

Forward Deployed AI Engineering — Capstone Project

Turning requirements into specifications and prompts. Your requirements describe what the system must do. This stage turns each of them into something precise enough to build against and to test.

Week Two · Days One and Two

Why the Prompt Library Exists

In a conventional software project, the step between a requirement and working code is a technical specification. In a system built around language models, a large part of that specification is the prompt itself. The prompt is not a throwaway string that lives in the middle of a function; it is the place where a requirement becomes behaviour, and it deserves the same version control and review as any other design artefact.

Students who keep their prompts scattered through the codebase end up unable to answer simple questions: Which requirement does this instruction serve? What did it say last week? Why was that example added? A library solves all three problems and costs very little to maintain if you start it on day one.

Versioning Policy

Every prompt change creates a new versioned file. The previous version is never overwritten.

Rule	Detail
File naming	<code>prompts/{name}_prompt_v{major}_{minor}.md</code> — e.g. <code>intent_classifier_prompt_v1_1.md</code> . The unversioned file (<code>intent_classifier_prompt.md</code>) always points to the current production version and is replaced by copying the new version over it.
Version numbers	Increment the minor version for wording tweaks and example additions (v1.0 → v1.1). Increment the major version for structural changes to steps, output schema, or routing logic (v1.x → v2.0).
Prompt library update	Update the register entry for the changed prompt: bump the version field, add a line to the change history, and update “known weaknesses”

Rule	Detail
	if the change was motivated by a failure mode.
Eval harness update	Update <code>SYSTEM_PROMPT</code> and <code>USER_PROMPT_TEMPLATE</code> in the corresponding <code>run*_eval.py</code> file to match the new prompt. Re-run with a new <code>--experiment-name</code> that includes the version (e.g. <code>intent-v1.1-gemini</code>).
Old versions	Kept in <code>prompts/</code> indefinitely. Never delete a versioned file — it is the record of what ran and what its results were.

The Four Categories

Category	What these prompts do	When you use them
Specification prompts	Convert a requirement into a technical specification and a set of acceptance criteria	At the start of week two, before you build
Build prompts	The instructions that run inside the system itself: classification, retrieval framing, answer drafting	Throughout the build
Review prompts	Critique your own specifications and code against the requirements they came from	After each component is drafted
Evaluation prompts	Judge answer quality, detect unsupported claims, compare against expert answers	During the evaluation run in week two

Section 1: From Requirement to Specification

The prompts in this library directly cover FR-01, FR-02, FR-03, FR-04, FR-05, and FR-11. FR-06 through FR-10 and FR-12 are operational or infrastructure concerns handled outside the prompt layer.

Requirement ID	Specification summary	Inputs	Outputs	Acceptance criteria
FR-01	Three classifiers run in sequence: fluency (PR-02), intent (PR-01), urgency (PR-03). Each produces a structured JSON output with a calibrated confidence score. Fluency output gates intent interpretation; intent	Raw ticket fields: ticket_id, channel, subject, body, customer_tier, received_at	language_fluency, intent, urgency, must_not_auto_respond, three confidence scores	Each ticket leaving classification has non-null values for all five fields. Confidence calibration within 5pp of observed precision per confidence band on validation set (NFR-03).

Requirement ID	Specification summary	Inputs	Outputs	Acceptance criteria
	output gates urgency and routing.			
FR-02	Intent classifier (PR-01) sets <code>must_not_auto_respond = true</code> for the six always-escalate intents. Coordinator (PR-00) notes FLAG A but continues the full pipeline through urgency, RAG, and <code>response_drafter</code> — draft and <code>relevant_doc_ids</code> are included in the escalation payload for the human reviewer. Only EXIT B (<code>answerable=false</code>) triggers an early stop.	intent_classifier output: intent, <code>must_not_auto_respond</code>	Route = escalate; escalation reason logged; draft and <code>relevant_doc_ids</code> included in payload for human reviewer	Zero tickets with <code>must_not_auto_respond = true</code> are auto-responded. MNR escalations carry a draft and <code>relevant_doc_ids</code> for the human reviewer. Verified against all 87 development tickets with always-escalate intents.
FR-03	RAG agent (PR-04) receives a pre-retrieved candidate set and filters it to the documents that genuinely address the ticket's core question. Assesses whether the filtered set is sufficient to draft a complete response.	ticket body, intent, urgency, candidate documents from vector search	<code>relevant_doc_ids</code> , <code>answerable (true/false)</code> , confidence	<code>doc_precision</code> ≥ 0.80 , <code>doc_recall</code> ≥ 0.80 on answerable tickets in development set. <code>answerable_accuracy</code> ≥ 0.85 .
FR-04	Response drafter (PR-05) writes a customer-facing response citing only the provided documents. Every factual claim must be attributable to a specific document. A post-draft hallucination check step reviews each sentence before output.	ticket fields, upstream classifications, relevant documents	draft text, citations list, <code>answered_fully</code> flag, confidence	<code>citation_coverage</code> ≥ 0.90 . <code>must_not_claim_rate</code> = 0.0. <code>must_mention_coverage</code> ≥ 0.85 . No claim in the draft is unsupported by a cited document.
FR-05	Coordinator (PR-00) collects all sub-agent outputs before making a	All sub-agent outputs accumulated in <code>completed_invocations</code>	<code>final_decision</code> with route, <code>escalation_reason</code> , confidence; upstream	100% of escalated tickets carry a <code>final_decision</code> with a

Requirement ID	Specification summary	Inputs	Outputs	Acceptance criteria
	final_decision. For escalated tickets, the final_decision payload includes the intent, confidence, relevant_doc_ids (if any), and escalation_reason — giving the human agent full context on arrival.		context available for context packet construction	populated escalation_reason.
FR-11	Intent classifier (PR-01, Step 1) restates the non-fluent ticket in clear English before extracting technical signals. Fluency classifier (PR-02) upstream flags non-fluent tickets so Step 1 knows to apply restatement.	language_fluency flag from PR-02, raw ticket body	Normalised interpretation used throughout classification; restatement visible in reasoning_summary	Non-fluent classification accuracy not materially below fluent accuracy on development set.

Section 2: The Prompt Register

Prompt PR-00

Field	Value
Name and purpose	Coordinator — orchestrates the full triage pipeline. Receives each raw ticket and the list of completed sub-agent invocations, then outputs exactly one action per turn: either invoke the next sub-agent with explicit inputs, or make a final routing decision.
Category	Build
Serves requirement	FR-02 (escalation routing), FR-05 (escalation context packet), FR-07 (decision log)
Version	1.1
Model used	google/gemini-3.1-flash-lite via OpenRouter

Field	Value
Inputs it expects	Ticket fields (ticket_id, channel, subject, body, customer_tier, received_at), plus the completed_invocations block (all prior agent IDs, inputs, and outputs)
Output format required	Single JSON object: <code>{next_action, agent_call, final_decision, reasoning}</code> . Exactly one of agent_call or final_decision is non-null.
How you know it worked	routing_accuracy ≥ 0.90 , escalation_safety ≥ 0.95 (no under-triage), pipeline_efficiency ≥ 0.80 (correct turn counts per exit path) on development set. Zero under-triage events (predicted auto_respond when ticket should escalate). MNR escalations include non-null draft and relevant_doc_ids.
Known weaknesses	Up to 6 sequential LLM calls per ticket — increases latency and cost. Relies on the model following the multi-step format precisely; any malformed turn wastes a call. EXIT B requires the model to correctly recognise answerable=false from RAG output.
Change history	v1.0 — initial; EXIT A (MNR) short-circuited pipeline after PR-01, skipping urgency/RAG/draft. v1.1 (2026-09-17) — EXIT A converted to FLAG A: MNR tickets now run the full pipeline; draft and relevant_doc_ids are included in the escalation payload for the human reviewer. Only EXIT B (answerable=false) triggers an early stop.

System prompt:

You are the coordinator agent for CloudServe Solutions' support triage pipeline. You orchestrate a set of specialist sub-agents to classify, assess, and respond to incoming support tickets.

You operate in a loop. On each turn you receive:

1. The original support ticket
2. A list of sub-agent invocations completed so far (agent ID, inputs given, output received)

On each turn you output exactly one action:

- invoke_agent: call the next sub-agent and specify its exact inputs
- final_decision: end the pipeline with a routing decision

No sub-agent communicates with another. You are the only one who reads their outputs and decides what to do next.

Your output must be a single JSON object. Do not include any text after the JSON object.

User prompt template:

You are coordinating the triage pipeline for the following support ticket.

--- TICKET ---

Ticket ID: {ticket_id}
Channel: {channel}
Subject: {subject}
Body: {body}
Customer Tier: {customer_tier}
Received At: {received_at}

--- END TICKET ---

--- COMPLETED INVOCATIONS ---

{completed_invocations_block}

--- END COMPLETED INVOCATIONS ---

Work through each step:

STEP 1 – REVIEW COMPLETED WORK

List what has been completed so far and what outputs are now available.

If no invocations have been completed, state "Pipeline not started."

STEP 2 – CHECK ROUTING FLAGS

Check the following flags – they affect the final route but do NOT stop the pipeline early:

FLAG A – must_not_auto_respond = true (from intent_classifier):

The intent class requires human review. Note this flag but DO NOT stop here.
Continue the pipeline normally through urgency, RAG, and response_drafter.
The draft will be included in the escalation payload for the human reviewer.
Make the final_decision (escalate, with draft) only after all agents have run.

EXIT B – Escalate immediately if:

- rag_agent output shows answerable = false
Reason: no documentation covers this ticket; the response_drafter cannot produce
a useful draft. Escalate after RAG – do not invoke response_drafter.

Only EXIT B triggers an early stop. FLAG A does not.

STEP 3 – DETERMINE THE NEXT ACTION

If EXIT B has not fired, invoke agents in this order:

1. fluency_classifier (no dependencies)
2. intent_classifier (requires: language_fluency from step 1)
3. urgency_classifier (requires: intent from step 2)
4. rag_agent (requires: intent and urgency from steps 2-3)
5. response_drafter (requires: all prior outputs – runs even if FLAG A is set)

Only invoke an agent if all its required inputs are available from completed invocations or from the original ticket.

STEP 4 – OUTPUT FINAL JSON

If invoking an agent:

{

```
"next_action": "invoke_agent",
"agent_call": {
  "agent_id": "<PR-XX>",
  "agent_name": "<name>",
  "inputs": { <exact key-value pairs the agent requires> }
},
"final_decision": null,
"reasoning": "<one sentence>"
}

If making a final decision:
{
  "next_action": "final_decision",
  "agent_call": null,
  "final_decision": {
    "route": "<auto_respond or escalate>",
    "escalation_reason": <null or one sentence explaining why>,
    "draft": <null or the draft string from response_drafter output>,
    "confidence": <float 0.0–1.0>
  },
  "reasoning": "<one sentence>"
}
```

Full prompt with examples, pipeline trace, and design rationale: [prompts/coordinator_prompt.md](#)

Prompt PR-01

Field	Value
Name and purpose	Intent classifier — classify each ticket into one of 22 intent classes; flag the six always-escalate intents; emit a calibrated confidence score.
Category	Build
Serves requirement	FR-01 (intent classification), FR-02 (must_not_auto_respond flag), FR-11 (non-fluent restatement in Step 1)
Version	1.1
Model used	<code>google/gemini-3.1-flash-lite</code> via OpenRouter
Inputs it expects	ticket_id, channel, subject, body, customer_tier, language_fluency
Output format required	<code>{"intent": string, "confidence": float, "must_not_auto_respond": boolean, "reasoning_summary": string}</code>

Field	Value
How you know it worked	intent accuracy ≥ 0.85 , escalation_safety ≥ 0.95 , confidence_on_wrong ≤ 0.65 on development set (eval: <code>run_langsmith_eval.py</code>)
Known weaknesses	Low-data regime (~22 examples per intent in dev set) — per-category precision varies widely. confidence_on_wrong currently above ≤ 0.65 target (parked for tuning). Non-fluent restatement in Step 1 is implicit, not enforced structurally. <code>account_access</code> vs <code>configuration_help</code> boundary is ambiguous for permission-management tickets.
Change history	v1.0 — initial. Six-step chain-of-thought: decode → extract signals → hard-escalate screen → select intent → calibrate confidence → output JSON. v1.1 (2026-09-16) — Tightened <code>security_incident</code> (401/rotation excluded); tightened <code>unclear_request</code> (threshold: <2 signals); added STEP 4 disambiguation guide for 4 confusion pairs (<code>api_key_issue</code> / <code>security_incident</code> , <code>quota_or_overage</code> / <code>billing_query</code> , <code>api_usage_question</code> / <code>data_export</code> , <code>authentication_failure</code> / <code>account_access</code>).

System prompt:

```
You are a support ticket classifier for CloudServe Solutions, a cloud infrastructure platform.
Your job is to read a support ticket and return a structured JSON classification.

You MUST work through exactly six steps before outputting your answer.
Do not skip steps. Do not merge steps. Show your reasoning for each step.

Your final output must be a single JSON object matching this schema:
{
  "intent": string,
  "confidence": float (0.0-1.0),
  "must_not_auto_respond": boolean,
  "reasoning_summary": string
}

Do not include any text after the JSON object.
```

User prompt template:

```
Classify the following support ticket by working through the six steps below.

--- TICKET ---
Ticket ID:      {ticket_id}
Channel:        {channel}
Subject:        {subject}
Body:           {body}
```

```
Customer Tier: {customer_tier}
Language:      {language_fluency}
--- END TICKET ---
```

--- INTENT CLASSES ---

```
account_access, api_key_issue, api_usage_question, authentication_failure,
billing_query, compliance_request, configuration_help, data_export,
data_residency, database_issue, deployment_failure, feature_request,
integration_help, onboarding, performance_degradation, quota_or_overage,
rate_limit, rollback_request, security_incident, sso_configuration,
unclear_request, webhook_issue
```

--- END INTENT CLASSES ---

Work through each step:

STEP 1 – DECODE THE REQUEST

Read the ticket carefully. If the language is non-fluent or abbreviated, restate the customer's actual problem in clear English in one sentence. Note any implied urgency cues (words like "urgent", "down", "breaking", "can't", production timelines, or revenue impact).

STEP 2 – EXTRACT TECHNICAL SIGNALS

List the specific technical terms, product areas, error codes, and action verbs present in the ticket. These are your primary classification features. Be precise – "401" is more specific than "error".

STEP 3 – SCREEN FOR HARD-ESCALATE INTENTS

Check whether the ticket matches any of these six always-escalate intents:

- security_incident: confirmed breach or active compromise – an unauthorized party accessed systems, data, or credentials, or there is direct evidence of ongoing unauthorized use. NOT security_incident: a key returning 401 errors or needing rotation is api_key_issue; secrets appearing in logs is configuration_help.
- compliance_request: audit records, regulatory documentation, data retention evidence
- feature_request: asking for a new capability, enhancement, or product change
- unclear_request: body provides fewer than two technical signals – no error codes, no product area, no action verb – making the actual problem impossible to determine. When uncertain between a vague technical intent and unclear_request, choose unclear_request.
- data_residency: questions about where data is stored, data sovereignty, regional data laws – always requires human review due to legal exposure
- billing_query: invoice disputes, charge questions, plan changes, payment issues – always requires human review due to financial exposure

If yes: state which one applies and mark must_not_auto_respond = true.

If no: state "No hard-escalate trigger found."

STEP 4 – SELECT THE PRIMARY INTENT

From the 22 intent classes, identify the single best match.

If two intents are competing, name both and explain in one sentence why one takes precedence over the other.

State your chosen intent and a brief justification.

Disambiguation guide for commonly confused pairs:

- `api_key_issue` vs `security_incident`: Choose `api_key_issue` when the ticket is about a key not working, needing rotation, or having wrong permissions – even if the key was accidentally exposed. Only choose `security_incident` if there is direct evidence an unauthorized party already used the credential to access systems or data.
- `quota_or_overage` vs `billing_query`: Choose `quota_or_overage` when the ticket is about hitting or exceeding usage limits, or overage charges due to plan limits. Choose `billing_query` for general invoice questions, payment issues, or charge explanations unrelated to plan quotas.
- `api_usage_question` vs `data_export`: Choose `api_usage_question` when the customer asks HOW to retrieve or paginate data via the API. Choose `data_export` only when the customer asks CloudServe to perform a bulk export on their behalf.
- `authentication_failure` vs `account_access`: Choose `authentication_failure` for any failed sign-in that is not an explicit account lock (OAuth, session, MFA, console login failures). Choose `account_access` only for explicit account locks or suspended accounts.

STEP 5 – CALIBRATE CONFIDENCE

Start at 0.80. Then adjust:

Raise to ≥ 0.85 only when ALL of the following are true:

- The ticket body contains multiple specific technical keywords (error codes, product names, action verbs) that point to exactly one intent class
- You did NOT consider any competing intent in Step 4
- The body is detailed enough that another reader would reach the same intent

Keep at 0.80 when:

- One clear intent with minor ambiguity (one competing intent briefly considered but clearly ruled out)

Lower to 0.65–0.79 when ANY of these apply:

- Two plausible intents competed and the choice required a judgment call
- Ticket body is short (1–2 sentences) or missing technical specifics
- Non-fluent language required interpretation to determine meaning
- Intent was inferred from context rather than explicit keywords

Lower to < 0.65 when TWO OR MORE lower conditions apply simultaneously, or when you genuinely cannot determine the dominant intent with confidence.

Justify your score in one sentence citing the specific signal(s) that drove it.

STEP 6 – OUTPUT FINAL JSON

Output the classification as a JSON object only. No prose after the object.

Set `must_not_auto_respond = true` ONLY if the intent is one of the six hard-escalate classes (`security_incident`, `compliance_request`, `feature_request`, `unclear_request`, `data_residency`, `billing_query`). Do not use confidence score to determine this flag.

```
{
  "intent": "<chosen intent>",
  "confidence": <float>,
  "must_not_auto_respond": <true|false>,
}
```

```
"reasoning_summary": "<one sentence summary of why this classification was chosen>"
}
```

Full prompt with worked example and design rationale: [prompts/intent_classifier_prompt.md](#)

Prompt PR-02

Field	Value
Name and purpose	Language fluency classifier — binary classification of each ticket's author as fluent or non_fluent English speaker. Runs first in the pipeline to inform downstream interpretation.
Category	Build
Serves requirement	FR-01 (language_fluency input to intent classifier), FR-11 (identifies tickets requiring restatement)
Version	1.0
Model used	google/gemini-3.1-flash-lite via OpenRouter
Inputs it expects	ticket_id, channel, body
Output format required	<code>{"fluency": "fluent" \ "non_fluent", "confidence": float, "reasoning_summary": string}</code>
How you know it worked	fluency_accuracy $\geq$ 0.90, non_fluent_recall $\geq$ 0.85 on development set (eval: run_fluency_eval.py). non_fluent_recall tracked separately due to class imbalance (120/500 = 24%).
Known weaknesses	Binary output — no “borderline” class for genuinely ambiguous cases. Typo-vs-non-fluency boundary is assessed by the model, not by a rule; individual borderline judgments may be inconsistent. Technical jargon and informality could be over-counted as non-fluency signals by an inattentive model.
Change history	v1.0 — initial. Three-step chain-of-thought: scan for signals → assess overall fluency → output JSON. Key rule: $\geq$ 2 distinct non-fluency signals required for non_fluent classification; typos alone are insufficient.

System prompt:

You are a language fluency detector for CloudServe Solutions, a cloud infrastructure platform. Your job is to read a support ticket body and classify the author's English fluency.

You MUST work through exactly three steps before outputting your answer.
Do not skip steps. Do not merge steps. Show your reasoning for each step.

Your final output must be a single JSON object matching this schema:

```
{
  "fluency": "fluent" | "non_fluent",
  "confidence": float (0.0-1.0),
  "reasoning_summary": string
}
```

Do not include any text after the JSON object.

User prompt template:

Classify the language fluency of the following support ticket body.

```
--- TICKET ---
Ticket ID: {ticket_id}
Channel:   {channel}
Body:      {body}
--- END TICKET ---
```

Work through each step:

STEP 1 – SCAN FOR NON-FLUENCY SIGNALS

Look for specific markers of non-native English writing:

- Subject-verb agreement errors ("builds that work last week are now fail")
- Missing or incorrect articles ("the", "a", "an")
- Unusual word order or sentence structure
- Non-standard verb tense or aspect ("we are having not change")
- Dropped pronouns or prepositions
- Literal translations that produce unnatural phrasing

List each signal you find, or state "No non-fluency signals found."

STEP 2 – ASSESS OVERALL FLUENCY

Consider the ticket as a whole:

- If you found two or more distinct non-fluency signals: classify as non_fluent
- If you found one borderline signal (could be a typo or autocorrect): weigh against the rest of the text
- If the text is grammatically natural, even if informal or abbreviated: classify as fluent

Note: technical jargon, abbreviations, and casual tone are NOT non-fluency signals.

Typos alone are NOT sufficient – all writers make typos.

STEP 3 – OUTPUT FINAL JSON

```
{
  "fluency": "<fluent or non_fluent>",
  "confidence": <float>,

```

```
"reasoning_summary": "<one sentence citing the specific signals>"
}
```

Full prompt with worked examples and design rationale: [prompts/language_fluency_prompt.md](https://github.com/GoogleCloudPlatform/ai-docs/blob/main/prompts/language_fluency_prompt.md)

Prompt PR-03

Field	Value
Name and purpose	Urgency classifier — classify each ticket's urgency as high, medium, or low. Runs after intent classification; urgency feeds both queue prioritisation and the response drafter's tone adaptation.
Category	Build
Serves requirement	FR-01 (urgency component of classification), FR-12 (urgency tier used in queue ordering)
Version	3.5
Model used	<code>google/gemini-3.1-flash-lite</code> via OpenRouter
Inputs it expects	<code>ticket_id</code> , <code>channel</code> , <code>subject</code> , <code>body</code> , <code>customer_tier</code> , <code>intent</code>
Output format required	<pre>{"urgency": "high" \ "medium" \ "low", "confidence": int (0-100), "reasoning_summary": string}</pre>
How you know it worked	<code>urgency_accuracy</code> $\geq$ 0.75, <code>urgency_safety</code> $\geq$ 0.90 (no under-triage), <code>high_urgency_recall</code> $\geq$ 0.90 on development set (eval: <code>run_urgency_eval.py</code>). Safety score: under-triage = 0.0, over-triage = 0.5, correct = 1.0.
Known weaknesses	Dataset label inconsistencies for <code>rollback_request</code> (identical bodies labeled HIGH vs MEDIUM) — unfixable via prompt. Credentials exposed in public repo classified HIGH by prompt but labeled MEDIUM in some dataset rows — intentional safety divergence. DEV-0026 (pagination data integrity, labeled HIGH) reads as how-to question in body. Over-triage is the accepted failure mode; under-triage is dangerous.
Change history	v1.0 — medium bias (accuracy = 0.452). v2.0 — low bias (accuracy = 0.392). v3.0 (2026-09-14) — OR-gate HIGH, AND-gate LOW, MEDIUM safe default. v3.1 (2026-09-16) — removed <code>rollback_request</code> auto-HIGH, added credential exposure, multi-user qualifier. v3.2 (2026-09-17) — added “deployed release causing errors” to HIGH, explicit NOT-LOW list. v3.3 (2026-09-17) — added service-component-stopped

Field	Value
	to HIGH, stuck batch jobs / always-escalate intents / departing employee to NOT-LOW; confidence changed to int scale. v3.4 (2026-09-17) — added auth failure and platform-blocking-operation — reverted due to over-triage. v3.5 (2026-09-17, current) — reverted v3.4 additions; retained rate limits hit in production as NOT-LOW.

System prompt:

You are an urgency classifier for CloudServe Solutions, a cloud infrastructure platform.
Your job is to read a support ticket and classify its urgency as high, medium, or low.

You MUST work through exactly four steps before outputting your answer.
Do not skip steps. Do not merge steps. Show your reasoning for each step.

Your final output must be a single JSON object matching this schema:

```
{
  "urgency": "high" | "medium" | "low",
  "confidence": int (0-100),
  "reasoning_summary": string
}
```

Do not include any text after the JSON object.

User prompt template:

Classify the urgency of the following support ticket.

--- TICKET ---

Ticket ID: {ticket_id}
Channel: {channel}
Subject: {subject}
Body: {body}
Customer Tier: {customer_tier}
Intent: {intent}

--- END TICKET ---

Work through each step:

STEP 1 – SCREEN FOR HIGH URGENCY

Answer YES if ANY of the following is true. If YES, urgency is HIGH – skip to Step 3.

- (a) Intent is security_incident
- (b) A production system is currently broken and customers or services cannot proceed:
 - all users locked out, service completely unreachable, build/deploy pipeline blocked with no workaround, a deployed release actively causing errors in production,
 - a service component or integration that has stopped functioning and not self-recovered (e.g. webhook delivery halted, log forwarding stopped overnight),
 - active data loss, active data exposure, or credentials exposed in a public location (e.g. public repo, public paste)

(c) Explicit urgency language refers to a live, broken production system affecting multiple users or services – not a single user's access issue:
"urgent", "ASAP", "down", "broken", "can't [access / deploy / connect]"

Answer NO if none of the above apply, or if the issue is in staging/development.
If NO → continue to Step 2.

STEP 2 – SCREEN FOR LOW URGENCY

Urgency is LOW only if BOTH of the following are true:

- (a) No active failure in production – issue is in staging/dev, or there is no failure at all
- (b) The request is informational or a configuration setup task with no production blocker:
a how-to question, initial setup/configuration of a new feature not yet in production,
a single-user non-critical access issue that does not block production work
(e.g. one developer cannot see a non-critical project)

NOT LOW – always MEDIUM:

- billing_query, data_residency, compliance_request, or unclear_request intents
(always-escalate categories that require human review)
- compliance, audit, or data-retention requests (implied deadline, requires action)
- feature requests (require product team triage, not purely informational)
- access revocation for a departing employee (time-sensitive security hygiene)
- access issues preventing production operations or affecting a team or service
- a batch job, export, or database restore that has been queued or running significantly longer than expected (possible pipeline failure)
- rate limits or quotas currently being hit in production (API calls returning errors)
- automated jobs failing or incorrect config actively causing production problems
- ongoing security concern: credentials, secrets, or sensitive data currently appearing in logs or outputs (even if framed as a how-to question)

If both (a) AND (b) are true → urgency is LOW.
If either is false → urgency is MEDIUM.

MEDIUM is the safe default for any operational problem not yet classified:
degraded but running, one of multiple things broken, automated jobs failing,
a compliance or audit request, a time-sensitive integration, ongoing security concerns
(e.g. sensitive data appearing in logs), a recent production incident requiring
a configuration change to prevent recurrence, or anything where the customer is
reporting a real problem but the service is still partially functional.

STEP 3 – CALIBRATE CONFIDENCE

Start at 80. Then adjust:

Raise to ≥ 85 if the dominant signal is explicit (exact keyword match, clear production context, or a definitive HIGH intent – security_incident).
Lower to 65-79 if urgency was inferred from context without explicit keywords.
Lower to < 65 if the body is vague or signals are mixed.

STEP 4 – OUTPUT FINAL JSON

```
{  
  "urgency": "<high, medium, or low>",  
  "confidence": <int>,  
}
```

```
"reasoning_summary": "<one sentence citing the dominant signal>"
}
```

Full prompt with worked examples, version history, and design rationale: [prompts/urgency_classifier_prompt.md](#)

Prompt PR-04

Field	Value
Name and purpose	RAG agent — given a support ticket and candidate documents pre-fetched via <code>src/retrieve.py</code> (ChromaDB + bidirectional graph expansion), identifies which documents are genuinely relevant and assesses whether the ticket can be fully answered from them.
Category	Build
Serves requirement	FR-03 (document retrieval relevance filtering), FR-05 (relevant_doc_ids in escalation context packet)
Version	3.0
Model used	<code>google/gemini-3.1-flash-lite</code> via OpenRouter
Inputs it expects	ticket_id, intent, urgency, subject, body, pre-fetched documents
Output format required	<code>{"relevant_doc_ids": [list], "answerable": boolean, "confidence": int (0-100), "reasoning_summary": string}</code>
How you know it worked	answerable_accuracy $\geq$ 0.85, doc_precision $\geq$ 0.80, doc_recall $\geq$ 0.80 on development set (eval: <code>run_rag_eval.py</code>). Precision/recall computed only for answerable tickets.
Known weaknesses	Strict answerability (“all steps present”) may produce conservative false negatives when one partial doc is available. Confidence score changed to integer scale 0–100.
Change history	v1.0 — initial. Four-step CoT: extract core question → assess docs → assess answerability → output JSON. Simulated retrieval. v2.0 — tool call retrieval via <code>retrieve_documents</code> . v3.0 (2026-09-17) — Retrieval moved to Python (<code>src/retrieve.py</code>); LLM receives pre-fetched docs. PARTIAL → SUPPORTING with explicit multi-symptom and subject/body divergence guidance. Answerability tightened: <code>true</code> only when ALL steps present; “when in doubt, prefer

Field	Value
	false". Confidence scale changed from float to int 0–100. Subject field added. Bidirectional graph expansion added.

System prompt:

You are a document relevance assessor for CloudServe Solutions, a cloud infrastructure platform.

You are given a support ticket and a set of candidate documents already retrieved from the knowledge base. Your job is to assess each document and decide whether the ticket is answerable.

BEFORE STEP 1 – identify distinct problems in this ticket:

Read the Subject and Body separately. List each distinct symptom or concern they describe.

Subject and Body often describe DIFFERENT aspects of the same case – both matter.

Each distinct problem you identify should be covered by at least one RELEVANT or SUPPORTING doc.

STEP 1 – ASSESS EACH DOCUMENT

For each document, state one of:

RELEVANT – directly addresses the core problem (same symptom, same product area, or contains the exact resolution steps needed)

SUPPORTING – covers a distinct symptom or related sub-topic described in the ticket that a thorough support response should address.

Mark SUPPORTING when:

- The ticket describes multiple distinct symptoms and this doc covers one of them
- The ticket subject and body point to different concerns – include a doc for each concern
- The doc covers a common co-occurring issue

Err on the side of inclusion: prefer SUPPORTING over NOT RELEVANT when any part of the ticket could reasonably reference this document.

NOT RELEVANT – covers a completely different domain

STEP 2 – ASSESS ANSWERABILITY

Based on RELEVANT + SUPPORTING documents together:

answerable = true ONLY if the documents provide ALL specific steps the customer needs – no key step missing, nothing invented

answerable = false if ANY of the following apply:

- No RELEVANT documents exist
- The ticket requires account-specific or live-system information not in any document
- The documents describe the symptom area but are missing the specific resolution procedure or commands needed
- The ticket asks for an action (rollback, restore, override) only a human can authorize

When in doubt, prefer answerable = false

STEP 3 – OUTPUT FINAL JSON

Output a single JSON object. Do not include any text after it.

```
{
  "relevant_doc_ids": [IDs of all RELEVANT and SUPPORTING docs; empty list if none],
  "answerable": true | false,
  "confidence": int (0–100),
  "reasoning_summary": string
}
```

```
}

confidence calibration:
  >= 85    explicit match – document covers the exact symptom or procedure described
  70–84    partial match – general area covered, not all specifics
  < 70     weak match – related but does not directly answer
```

User prompt template:

Assess the following support ticket against the retrieved documents below.

```
--- TICKET ---
Ticket ID: {ticket_id}
Intent:    {intent}
Urgency:   {urgency}
Subject:   {subject}
Body:      {body}
--- END TICKET ---
```

```
--- RETRIEVED DOCUMENTS ---
{documents}
--- END DOCUMENTS ---
```

Work through Steps 1–3 and output the JSON.

Full prompt with worked examples and design rationale: [prompts/rag_agent_prompt.md](#)

Prompt PR-05

Field	Value
Name and purpose	Response drafter — write a complete, attributed customer response using only the documents identified as relevant by the RAG agent. No claim may appear in the draft that cannot be traced to a specific provided document.
Category	Build
Serves requirement	FR-04 (grounded draft with citations), FR-02 (hard stop for must_not_auto_respond tickets), FR-06 (draft for agent review queue)
Version	1.0
Model used	<code>google/gemini-3.1-flash-lite</code> via OpenRouter
Inputs it expects	ticket_id, channel, body, customer_tier, language_fluency, intent, urgency, must_not_auto_respond, relevant documents

Field	Value
	(same block format as PR-04)
Output format required	<code>{"draft": string, "citations": [list], "answered_fully": boolean, "confidence": float, "reasoning_summary": string}</code>
How you know it worked	must_not_claim_rate = 0.0 (compliance — any value above 0 is a blocker), citation_coverage ≥ 0.90, must_mention_coverage ≥ 0.85, answered_fully_accuracy ≥ 0.80 on ground_truth_responses.json (eval: <code>run_response_draft_eval.py</code>).
Known weaknesses	Hallucination check is post-hoc (Step 4 reviews the completed draft) — not a preventive constraint during generation; a highly capable model may still rationalise unsupported claims as supported. Citations list doc IDs, not section identifiers — auditability is at article level, not paragraph level. must_not_claim enforcement relies on training signal rather than structural output schema.
Change history	v1.0 — initial. Five-step chain-of-thought: routing gate → map documents to question → draft → hallucination check (sentence-by-sentence attribution) → output JSON. Hard stop in Step 1 for must_not_auto_respond before any drafting begins. Language adaptation (fluent/non_fluent) and urgency adaptation (high/medium/low) applied in Step 3.

System prompt:

You are a support response drafter for CloudServe Solutions, a cloud infrastructure platform. Your job is to write a clear, accurate response to a customer support ticket using ONLY the provided documentation. You must not invent any information.

You MUST work through exactly five steps before outputting your answer. Do not skip steps. Do not merge steps. Show your reasoning for each step.

Your final output must be a single JSON object matching this schema:

```
{
  "draft": string,
  "citations": [list of doc IDs],
  "answered_fully": true | false,
  "confidence": float (0.0–1.0),
  "reasoning_summary": string
}
```

Do not include any text after the JSON object.

User prompt template:

Draft a response to the following support ticket using only the provided documents.

--- TICKET ---

Ticket ID: {ticket_id}
Channel: {channel}
Body: {body}
Customer Tier: {customer_tier}
Language: {language_fluency}

--- END TICKET ---

--- UPSTREAM CLASSIFICATION ---

Intent: {intent}
Urgency: {urgency}
Must Not Auto-Respond: {must_not_auto_respond}

--- END UPSTREAM CLASSIFICATION ---

--- RELEVANT DOCUMENTS ---

{documents_block}

--- END RELEVANT DOCUMENTS ---

Work through each step:

STEP 1 – CHECK ROUTING GATE

If must_not_auto_respond is true, stop here.

Output a brief internal escalation placeholder:

"This ticket has been flagged for human review and will not receive an automated response."

Set answered_fully = false, confidence = 1.0, citations = [].

Skip to Step 5.

STEP 2 – MAP DOCUMENTS TO THE CUSTOMER'S QUESTION

Identify which parts of the customer's question each document addresses.

Note any part of the question that is NOT covered by any document – you will need to acknowledge this gap explicitly in the draft.

STEP 3 – DRAFT THE RESPONSE

Write the customer-facing response following these rules:

- Open with a brief acknowledgement of the problem (one sentence)
- Present the resolution steps or information drawn directly from the documents
- For each factual claim, note internally which document it comes from (you will cite these in the output, not in the response text itself)
- If a part of the question cannot be answered from the documents, include a sentence such as: "I am not able to confirm [X] from our documentation – a member of our team will follow up on this point."
- Close with a next step or offer of further assistance

Language adaptation:

- fluent: standard professional tone, technical terms acceptable
- non_fluent: short sentences, plain vocabulary, avoid idioms and abbreviations

Urgency adaptation:

- high: action-first, skip pleasantries, lead with the most important step
- medium: professional and clear

- low: conversational, can include brief context before the resolution

STEP 4 – VERIFY NO HALLUCINATION

Review your draft sentence by sentence.

For each factual claim, confirm it appears explicitly in one of the provided documents.

If you find any sentence that you cannot attribute to a specific document:

- Either remove it, or
- Replace it with an explicit acknowledgement of uncertainty

List the doc IDs that provided content for the draft.

STEP 5 – OUTPUT FINAL JSON

```
{  
  "draft": "<the complete response text>",  
  "citations": [<list of doc IDs that contributed content>],  
  "answered_fully": <true if every aspect addressed, false if any gap acknowledged>,  
  "confidence": <float>,  
  "reasoning_summary": "<one sentence on coverage and any gaps>"  
}
```

Full prompt with worked examples and design rationale: [prompts/response_draft_prompt.md](#)

Section 3: What Makes a Prompt Worth Keeping

Check	What you are looking for	Status across this library
Does it state the role and the task separately?	A prompt that mixes who the model is with what it should do tends to produce inconsistent output.	All prompts open the system prompt with a role sentence and a separate task sentence.
Are the inputs clearly delimited?	Ticket text and instructions must be distinguishable, or a customer can accidentally instruct your system.	All prompts wrap ticket content in <code>--- TICKET ---</code> / <code>--- END TICKET ---</code> blocks, separating it from instruction text. Injection is partially mitigated — ticket body is in a named block, not interleaved with instructions.
Is the output format specified exactly?	If you are parsing the response, the format is part of the contract and must be stated.	Each prompt specifies exact JSON field names, types, and value constraints in both the system prompt and the output step.
Does it say what to do when the	Without this instruction the model will invent something.	PR-04 defines answerable = false conditions explicitly. PR-05 requires explicit gap acknowledgement in the draft. PR-01 has <code>unclear_request</code> for undeterminable intent.

Check	What you are looking for	Status across this library
answer is not known?		
Are the examples representative?	Examples drawn only from easy cases teach the model that every case is easy.	PR-01, PR-03, PR-04, PR-05 include both positive and negative examples (e.g. escalation path, partial coverage, non-fluent input). PR-02 includes a borderline fluent case.
Does it forbid what must never happen?	Constraints stated positively are followed more reliably than constraints implied by omission.	PR-05 Step 1 is a hard stop before drafting for <code>must_not_auto_respond</code> . PR-04 Step 3 explicitly enumerates when answerable = false. PR-01 Step 3 lists the six always-escalate intents with trigger phrases.

Injection handling note: The ticket body field in all user prompts is placed inside a clearly labelled `--- TICKET ---` block that follows the classification instructions. A customer cannot inject instructions that appear before or inside the instruction set. The remaining risk is jailbreak-style injection within the body field itself — this is not structurally mitigated in the current prompt design and should be tested against the synthetic PII/injection test set (FR-09).

Section 4: Traceability Check

Requirement ID	Specification written?	Prompts covering it	Evaluation harness	Gaps
FR-01	Yes — Section 1 above	PR-01 (intent, confidence), PR-02 (fluency), PR-03 (urgency)	<code>run_langsmith_eval.py</code> , <code>run_fluency_eval.py</code> , <code>run_urgency_eval.py</code>	PR-01 confidence_on_wrong above ≤ 0.65 target — parked for tuning
FR-02	Yes — Section 1 above	PR-01 (<code>must_not_auto_respond</code> flag), PR-00 (FLAG A routing — full pipeline runs, final route=escalate)	<code>run_langsmith_eval.py</code> (escalation_safety), <code>run_coordinator_eval.py</code> (routing_accuracy)	None — MNR escalations now include draft + relevant_doc_ids for human reviewer
FR-03	Yes — Section 1 above	PR-04 (relevance filtering + answerability)	<code>run_rag_eval.py</code>	Retrieval is simulated in eval — real vector search recall not yet measured
FR-04	Yes — Section 1 above	PR-05 (grounded draft with citations)	<code>run_response_draft_eval.py</code>	<code>must_not_claim</code> enforcement relies on model compliance, not structural output constraint

Requirement ID	Specification written?	Prompts covering it	Evaluation harness	Gaps
FR-05	Yes — Section 1 above	PR-00 (escalation_reason in final_decision), PR-04 (relevant_doc_ids forwarded)	run_coordinator_eval.py	Context packet schema (FR-05) is assembled at orchestration layer — not part of a single prompt
FR-06	Yes — PR-00 v1.1	PR-00 (FLAG A path — draft included in escalation payload), PR-05 (draft produced for all answerable tickets)	run_coordinator_eval.py	MNR escalations now carry draft + relevant_doc_ids; human reviewer gets a complete work package
FR-07	No prompt required	—	—	Decision log is written by the orchestration layer from coordinator outputs
FR-08	No prompt required	—	—	Channel normalisation is handled at ingest before any prompt is called
FR-09	Not yet covered	—	—	PII validation guardrail prompt not yet designed — required before production deployment
FR-10	No prompt required	—	—	Fairness metrics are computed from eval results, not produced by a prompt
FR-11	Yes — Section 1 above	PR-02 (flags non-fluent), PR-01 Step 1 (restatement)	run_fluency_eval.py (non_fluent_recall), run_langsmith_eval.py (per-fluency accuracy)	Restatement in PR-01 Step 1 is implicit — not a structured output field
FR-12	No prompt required	—	—	Queue ordering is an application layer concern

Prompt SYS-00 (CLAUDE.md)

Field	Value
Name and purpose	Developer tooling system prompt — provides Claude Code (claude.ai/code) with project context, environment setup, architecture overview, prompt versioning policy, and communication style constraints. Functions as a persistent system prompt injected into every AI-assisted development session on this codebase.
Category	Build
Serves requirement	All requirements — governs how the AI engineering assistant approaches this codebase during development, evaluation, and prompt tuning
Version	1.0
Location	<code>.claude/CLAUDE.md</code> (auto-loaded by Claude Code at session start)
Inputs it expects	N/A — loaded automatically; not called with runtime inputs
Output format required	N/A — controls assistant behaviour, not an output schema
How you know it worked	Claude Code follows the prompt versioning policy without reminder; does not give terminal commands in chat; reads requirements before assuming scope; correctly identifies FLAG A vs EXIT B routing.
Known weaknesses	Architecture diagram still shows <code>EXIT A</code> (stale — superseded by FLAG A in PR-00 v1.1). Env-var guidance still names <code>OPENROUTER_API_KEY</code> as required; Groq and Gemini keys are now supported and take priority.
Change history	v1.0 — initial. Contains environment setup, eval commands, architecture overview with EXIT A, metric targets, prompt versioning policy, requirements guidance, and communication style rules.

Full content:

```
# CLAUDE.md

This file provides guidance to Claude Code (claude.ai/code) when working with code in this repository.

## Environment

Python 3.14 managed by `uv`. Installed packages: `openai>=1.0.0`, `python-dotenv`, `chromadb`.

```bash
uv pip install "openai>=1.0.0" python-dotenv chromadb
```

Env vars are auto-loaded from `.env` at the project root. Required: `OPENROUTER_API_KEY` . Optional: `MODEL_NAME` (default: `google/gemini-3.1-flash-lite` ).

## Running Evaluations

All eval commands run from the `evaluation/` directory. Common args:

Arg	Purpose
<code>--sample N</code>	First N tickets only (smoke test)
<code>--experiment-name STR</code>	CSV output name; auto-timestamped if omitted
<code>--concurrency N</code>	Parallel LLM calls (default 4; reduce to 2 on rate limits)

```
python run_langsmith_eval.py --sample 20 --experiment-name "intent-smoke"
python run_fluency_eval.py --sample 20 --experiment-name "fluency-smoke"
python run_urgency_eval.py --sample 20 --experiment-name "urgency-smoke"
python run_rag_eval.py --sample 20 --experiment-name "rag-smoke"
python run_response_draft_eval.py --sample 20 --experiment-name "draft-smoke"
python run_coordinator_eval.py --sample 20 --experiment-name "coord-smoke"
```

Results → `evaluation/results/{experiment_name}.csv` .

`run_response_draft_eval.py` uses `ground_truth_responses.json` (200 tickets). All others use `development_tickets.json` (500 tickets). Both are under `requirements/05_Datasets/` .

## Architecture

The system is a coordinator-mediated multi-agent triage pipeline. No application server exists yet — only prompts and evaluation harnesses.

```
PR-00 Coordinator (loop — one action per turn)
├─ PR-02 fluency_classifier → language_fluency
├─ PR-01 intent_classifier → intent, must_not_auto_respond → EXIT A (escalate)
├─ PR-03 urgency_classifier → urgency
├─ PR-04 rag_agent → relevant_doc_ids, answerable → EXIT B (escalate)
└─ PR-05 response_drafter → draft, citations → auto_respond
```

- **Coordinator (PR-00)** is called in a loop; each turn outputs exactly one of `invoke_agent` or `final_decision` . Sub-agents never communicate with each other — all data flows through the coordinator.
- **EXIT A:** `must_not_auto_respond = true` after PR-01 → escalate without calling PR-03/04/05.
- **EXIT B:** `answerable = false` after PR-04 → escalate without calling PR-05.
- Each **eval harness embeds its own copy** of `SYSTEM_PROMPT` and `USER_PROMPT_TEMPLATE` . Changing a prompt requires updating both the `prompts/*.md` file and the harness.
- The **coordinator harness mocks sub-agents** via `get_mock_output()` — no real sub-agent LLM calls.

- The **response draft harness uses perfect RAG** — passes `expected_doc_ids` directly to isolate draft quality from retrieval quality.

### Metric targets

Agent	Critical metric	Target
PR-02 fluency	<code>non_fluent_recall</code>	≥ 0.85
PR-01 intent	<code>escalation_safety</code>	≥ 0.95 · <code>confidence_on_wrong</code> ≤ 0.65 (parked)
PR-03 urgency	<code>urgency_safety</code> , <code>high_urgency_recall</code>	≥ 0.90
PR-04 RAG	<code>answerable_accuracy</code> / <code>doc_precision</code> / <code>doc_recall</code>	≥ 0.85 / 0.80 / 0.80
PR-05 draft	<code>must_not_claim_rate</code>	<b>0.0</b> — any value above 0 is a compliance blocker
PR-00 coordinator	<code>routing_accuracy</code> , zero under-triage	≥ 0.90

### Prompt Versioning Policy

Every prompt change must follow these steps — do not skip any:

1. Save the new version as `prompts/{name}_prompt_v{major}_{minor}.md`
2. Copy the new version over the unversioned live file ( `prompts/{name}_prompt.md` )
3. Update `requirements-md/02_Stage_Workbooks/Stage_3_Prompt_Library.md` — bump version field, add change history line, update known weaknesses
4. Sync `SYSTEM_PROMPT` and `USER_PROMPT_TEMPLATE` in `evaluation/run*_eval.py` ; re-run with a versioned experiment name (e.g. `intent-v1.1-gemini` )
5. Never delete versioned files — they are the audit trail

Version numbering: minor bump (v1.0 → v1.1) for wording/examples; major bump (v1.x → v2.0) for structural step or schema changes.

### Requirements

When in doubt about any requirement, business rule, or scope decision, read the files in `requirements/01_Read_First/` before proceeding and ask clarifying questions rather than assuming.

### Communication Style

- Never give terminal commands in chat — update the relevant guide or doc file instead ``

Related documents: [01\\_Discovery Workbook.pdf](#) · [02\\_Requirements PRD.pdf](#) · [04\\_Sprint Plan.pdf](#) · [05\\_PRD Revision Log.pdf](#) · [AyushYadav Effort Log.pdf](#)