

# Stage 5: PRD Revision Log

---

## Stage Five: PRD Revision Log

---

### Forward Deployed AI Engineering — Capstone Project

*Recording what your requirements got wrong and why. You are required to revise your requirements document at least once during the build. This is where you record what changed and what prompted the change.*

### Week Two into Week Three · Compulsory

---

### Why This Is Compulsory

---

A requirements document written before any code exists is a set of educated guesses. Some of those guesses will be wrong, and building is how you find out which ones. A student that reaches the end of three weeks with an unchanged PRD has either been extraordinarily lucky or has stopped consulting the document, and the second is far more common than the first.

This log is not a confession. It is evidence that you were paying attention. The students who score well here are the ones who can point to a specific thing they discovered while building, explain how it contradicted an assumption they had written down, and show what they changed as a result.

The revision sits in the middle of the build, not at the end.

---

### Section 1: Revision Summary

---

| Field                          | Value                                                  |
|--------------------------------|--------------------------------------------------------|
| New version number             | 2.0                                                    |
| Date of revision               | 2026-09-18 (Updated Sep 19)                            |
| Who carried out the revision   | Ayush Yadav (Forward Deployed AI Engineer)             |
| Who reviewed and agreed it     | Marcus Adeyemi (Sponsor) & Daniel Okonkwo (L2/L3 Lead) |
| Number of requirements changed | 5                                                      |
| Number of requirements added   | 3                                                      |

| Field                          | Value |
|--------------------------------|-------|
| Number of requirements removed | 0     |

---

## Section 2: The Changes

---

| Requirement ID                                       | What it said in version one                                                                                                                                                                                                                          | What it says now                                                                                                                                                                                                                | What prompted the change                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                             | Agreed by                       |
|------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------|
| <b>FR-01</b><br>(Classification Engine Architecture) | Train custom lightweight supervised ML classifiers (e.g. DeBERTa/Scikit-Learn) on historical ticket data for Intent and Urgency classification to minimize latency.                                                                                  | Use <b>few-shot LLM prompts</b> ( <code>src/agents/intent.py</code> & <code>src/agents/urgency.py</code> ) for V1 classification; define <b>fine-tuned ML models</b> as the long-term production Northstar once dataset scales. | The historical baseline was <b>100% manual human triage</b> , resulting in subjective human label discrepancies (identical system errors like <code>500 Internal Server Error</code> inconsistently labeled <code>medium</code> by one rep and <code>high</code> by another). Combined with extreme data sparsity across <b>22 intent classes</b> (~22 samples/class in the 500-ticket set), supervised ML models overfitted to human noise. Few-shot LLM prompting allowed enforcing objective, rule-aligned domain criteria in V1. | Marcus Adeyemi & Daniel Okonkwo |
| <b>FR-02</b><br>(Escalation & Pre-Drafting)          | Route must-not-auto tickets ( <code>security_incident</code> , <code>compliance_request</code> , <code>feature_request</code> , <code>unclear_request</code> ) directly to human queue with <b>no automated draft generated</b> (EXIT A early stop). | Converted EXIT A to <b>FLAG A</b> : Must-not-auto tickets run the full pipeline through RAG and Response Drafter. Pre-drafted suggested responses + cited canonical doc IDs are                                                 | Initial build testing showed early aborts (EXIT A) forced human agents right back into 100% manual drafting (~12                                                                                                                                                                                                                                                                                                                                                                                                                     | Marcus Adeyemi & Daniel Okonkwo |

| Requirement ID                              | What it said in version one                                                                                                          | What it says now                                                                                                                                                                                                                                    | What prompted the change                                                                                                                                                                                                                            | Agreed by                   |
|---------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------|
|                                             |                                                                                                                                      | embedded in the escalation context packet for human 1-click review.                                                                                                                                                                                 | min/ticket).<br>72.3% of escalations had null/placeholder drafts.<br>Switching to Copilot pre-drafting eliminated placeholders 100% and provided 91.9% pre-draft coverage, saving ~19 hours of human drafting effort.                               |                             |
| <b>FR-04</b><br>(Grounding & Validation)    | Strict sentence-by-sentence exact substring match against retrieved documents<br>( <code>extracted_summary_sentences_match</code> ). | Relaxed sentence-level strict blocking; enforced RAG candidate filtering ( <code>doc_precision</code> / <code>doc_recall</code> ) + Response Drafter sentence attribution + PII / toxic egress validation.                                          | Strict sentence matching caused massive false-positive blocks on valid drafts that rephrased or summarized documentation, dropping auto-respond rate to 33.8%.<br>Relaxing strict sentence matching raised FCR to 61.0% with 0 hallucinated claims. | Ines Varga & Sofia Restrepo |
| <b>FR-05</b><br>(Escalation Context Packet) | Context packet contains intent, confidence, doc IDs, and uncertainty statement.                                                      | Expanded context packet schema to include:<br><code>ticket_id</code> , <code>intent</code> , <code>urgency</code> , <code>confidence_score</code> , <code>relevant_doc_ids</code> , <code>suggested_draft</code> , <code>escalation_reason</code> , | L2/L3 support engineers requested full visibility into sub-agent reasoning, pre-drafted responses, and guardrail                                                                                                                                    | Daniel Okonkwo              |

| Requirement ID                                     | What it said in version one                                    | What it says now                                                                                                                                                                                                                                                                                    | What prompted the change                                                                                                                                              | Agreed by        |
|----------------------------------------------------|----------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------|
|                                                    |                                                                | <code>guardrail_results</code> , and <code>decision_id</code> .                                                                                                                                                                                                                                     | validation logs to make 1-click approval decisions (~1–2 min).                                                                                                        |                  |
| <b>FR-07</b><br>(Decision Logging & Governance)    | Log decisions with basic ticket attributes and routing action. | Log full structured decision payloads to SQLite database ( <code>storage/decisions.db</code> ), persisting all completed sub-agent invocations, inputs, outputs, guardrail results, and unique UUID decision IDs.                                                                                   | Required for 100% auditability during autumn compliance review and Prometheus telemetry monitoring.                                                                   | Marcus Adeyemi   |
| <b>FR-13 [NEW]</b><br>(OOP Domain Architecture)    | Unstructured script functions in codebase.                     | Codebase structured into modular Object-Oriented Programming (OOP) classes following SOLID principles ( <code>BaseAgent</code> , <code>DocumentRetriever</code> , <code>GuardrailPipeline</code> , <code>DecisionRepository</code> , <code>TicketNormalizer</code> , <code>MetricsService</code> ). | Required for maintainability, test isolation, and clean software architecture.                                                                                        | Engineering Lead |
| <b>FR-14 [NEW]</b><br>(3-Tier Operational Routing) | Legacy 2-tier routing (Auto-respond vs Escalation).            | 3-Tier Operational Routing: L1 Autonomous AI (61.0% FCR), L2 Human Copilot Approver (22.8% FCR), L3 Senior Specialist (16.2% Queue Load).                                                                                                                                                           | Reflects actual operational reality: 83.8% Total Effective FCR achieved by combining autonomous AI and 1-click human copilot approvals over the 100% manual baseline. | Marcus Adeyemi   |

| Requirement ID                                           | What it said in version one      | What it says now                                                                                                                                                | What prompted the change                                                                                                                                                                                 | Agreed by      |
|----------------------------------------------------------|----------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------|
| <b>FR-15 [NEW]</b><br>(Northstar Classification Roadmap) | Single-phase model architecture. | Phased classification roadmap: Phase 1 (LLM Few-Shot Prompting on small dataset) → Phase 2 (Fine-tuned DeBERTa-v3 ML Classifiers on >5,000 production tickets). | Transitioning to ML classifiers in Phase 2 guarantees deterministic inference latency (<50ms vs 1.5s LLM calls), eliminates API token costs for high-volume classification, and enhances predictability. | Marcus Adeyemi |

### Section 3: Assumptions That Turned Out to Be Wrong

| Assumption from version one                                                                                                                          | Did it hold? | What you found instead                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                    | What you changed because of it                                                                                                                                                                                                                                                                                                                                                                                           |
|------------------------------------------------------------------------------------------------------------------------------------------------------|--------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>1. Supervised ML models (e.g. SVM / Random Forest / DeBERTa) should be trained on the 500 dev tickets for intent &amp; urgency classification</b> | <b>NO</b>    | The pre-project baseline was <b>100% manual human triage</b> . Historical ticket labels exhibited <b>severe human subjective discrepancy and noise</b> (identical system errors like <code>500 Internal Server Error</code> were inconsistently labeled <code>medium</code> by one agent and <code>high</code> by another). Spreading 500 tickets across <b>22 intent classes</b> and 3 urgency levels resulted in extreme data sparsity (<8 samples per intent-urgency bucket). Supervised ML models overfitted to human | <b>Forced to use LLM few-shot prompting for V1:</b> Enforced clear, rule-aligned domain criteria in LLM system prompts ( <code>urgency_prompts.py</code> ) to overcome human label noise and data sparsity. Documented fine-tuned ML classifiers (DeBERTa-v3/SetFit) as the <b>production Northstar</b> once >5,000 clean production AI logs are recorded, reducing LLM costs and ensuring deterministic predictability. |

| Assumption from version one                                                                                                    | Did it hold? | What you found instead                                                                                                                                                                                                                | What you changed because of it                                                                                                                                                                                             |
|--------------------------------------------------------------------------------------------------------------------------------|--------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|                                                                                                                                |              | noise and failed to generalize.                                                                                                                                                                                                       |                                                                                                                                                                                                                            |
| <b>2. Must-Not-Auto tickets should abort early without drafting (EXIT A)</b>                                                   | <b>NO</b>    | Aborting early produced zero draft context for human agents, forcing them back into manual drafting (~12 min/ticket). 72.3% of escalations had null/placeholder drafts.                                                               | Changed to <b>FLAG A</b> : pipeline runs all agents and embeds pre-drafted suggested responses + cited doc IDs in the escalation packet for 1-click human approval (~1–2 min), saving ~19 hours of manual drafting effort. |
| <b>3. Sentence-level exact substring match is required for grounding</b><br>( <code>extracted_summary_sentences_match</code> ) | <b>NO</b>    | Strict sentence substring matching caused massive false-positive blocks on valid drafts that synthesized multiple KB articles, capping auto-respond rate at 33.8%.                                                                    | Relaxed strict sentence substring matching. Relied on RAG relevance filtering + Response Drafter attribution + PII/toxic egress guardrails, achieving 61.0% FCR with 0 false claims.                                       |
| <b>4. Non-fluent English tickets receive significantly lower customer CSAT scores</b>                                          | <b>NO</b>    | Analysis of 500 development tickets revealed non-fluent CSAT is <b>3.04 vs 2.95 for fluent speakers</b> (slightly higher!). The real penalty was <b>+14% resolution time</b> (466 min vs 408 min) due to human agent decoding delays. | Shifted focus of <code>FluencyClassifier</code> (PR-02) and <code>IntentClassifier</code> (PR-01 Step 1) to non-native query restatement to solve decoding delay (+14% time), rather than CSAT-focused routing.            |
| <b>5. Enterprise tier customers receive faster resolution in ~1 hour</b>                                                       | <b>NO</b>    | Enterprise tier actually had the <b>worst operational metrics</b> under manual support: 484 min resolution, 37.3% FCR, and 62.7% escalation rate due to higher technical complexity.                                                  | Focused primary automation throughput on Standard tier (51% of volume) while equipping Enterprise escalations with rich pre-drafted context packets.                                                                       |

## Section 4: What You Decided Not to Change

| What looked wrong                                                                           | Why you left it                                                                                                                                                                                                                                                                                                                               | What it would have cost to change                                                                                                                                                                                          | Revisit when?                                                                                                                                                            |
|---------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>LLM-based Intent &amp; Urgency Classification in V1 (despite higher API token costs)</b> | Kept LLM classification in V1 to handle data sparsity (500 tickets across 22 intents) and historical human label noise.                                                                                                                                                                                                                       | Attempting to train ML models on 500 noisy tickets caused accuracy to drop below 60% and introduced unpredictable misclassifications. LLM zero/few-shot prompting achieved 90.8% intent accuracy and 85.0% urgency safety. | Revisit in Phase 2 (Northstar) after collecting 5,000+ audited production AI logs to train fine-tuned DeBERTa-v3 ML classifiers.                                         |
| <b>Direct CSAT Upliftment Interventions (Categorized in MOSCoW “Won’t Have”)</b>            | Decided NOT to build prompt intervention loops or sentiment manipulation features for CSAT upliftment. Discovery proved CSAT (3.2/5) is an operational symptom of <b>response delays (8-12 hrs)</b> and <b>unresolved tickets (&lt;50% FCR)</b> , with an operational FCR ceiling of 3.41. CSAT cannot be inflated via prompt phrasing alone. | Attempting to manipulate CSAT via prompt sentiment smoothing adds prompt complexity without addressing response latency or documentation gaps.                                                                             | Revisit after FCR exceeds 60%+ and response times drop under 5 minutes, requiring product management to address non-pipeline drivers (doc quality, pricing, onboarding). |
| <b>billing_query and data_residency over-escalation (0% auto-respond in initial run)</b>    | Kept strict policy escalation for billing and data residency tickets in production baseline.                                                                                                                                                                                                                                                  | High legal and financial compliance exposure if an automated system issued an incorrect billing refund or regional data promise.                                                                                           | Revisit in v2.1 after 3 months of production L2 Copilot drafting logs demonstrate 99%+ human approval accuracy.                                                          |
| <b>Urgency classifier medium-bias (48.6% accuracy against ground truth)</b>                 | Left medium as the safe default for ambiguous operational tickets.                                                                                                                                                                                                                                                                            | Ground truth human labels contained subjective inconsistency (e.g. reps labeling identical outages as medium vs high). Forcing aggressive high/low classification risked under-triaging critical production outages.       | Revisit when fine-tuned ML classification models with audited labels are deployed in Phase 2.                                                                            |

## Section 5: The Reflection

| Question                                                                      | Your answer                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |
|-------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| What did you most misunderstand about the problem when you wrote version one? | I assumed that escalation was purely a binary “hand-off” where the AI stops completely, and that classical ML classifiers could be easily trained for intent/urgency. In reality, CloudServe’s pre-project workflow was <b>100% manual human triage</b> , so historical data contained significant human label noise. With only 500 tickets across 22 intent classes, supervised ML models failed, forcing us to use LLM few-shot prompting in V1 while targeting ML models as the long-term Northstar. Additionally, escalated tickets represent <b>56.2% of workload</b> , requiring an <b>AI Copilot</b> rather than an early exit. |
| Which piece of discovery work would have caught it earlier?                   | A detailed statistical audit of human label consistency across historical tickets. Noticing that different support reps labeled identical <b>500 Server Error</b> tickets as <b>medium</b> vs <b>high</b> would have immediately highlighted human label discrepancy and warned against training supervised ML classifiers on small, noisy datasets.                                                                                                                                                                                                                                                                                   |
| What would you do differently if you started this project again on Monday?    | I would build the 3-tier Copilot architecture (L1 Autonomous AI, L2 Copilot Approver, L3 Senior Specialist) from Day 1, and explicitly document the two-phase classification roadmap (V1 LLM few-shot prompting → V2 Northstar ML classifiers).                                                                                                                                                                                                                                                                                                                                                                                        |
| What is still uncertain, and what would you need to resolve it?               | The exact timeline to collect 5,000+ clean, audited production AI logs required to train the Northstar DeBERTa-v3 ML classifiers. Resolving this requires monitoring ticket volume growth over 3–6 months of production operation.                                                                                                                                                                                                                                                                                                                                                                                                     |

## Section 6: PRD Compliance Audit (v2.0 Verified)

| Requirement ID | Requirement Promise                                                                                                  | Implementation Status    | Empirical Evidence / Verification                                                                                              |
|----------------|----------------------------------------------------------------------------------------------------------------------|--------------------------|--------------------------------------------------------------------------------------------------------------------------------|
| FR-01          | Multi-agent intent & urgency classification with calibrated confidence.                                              | <b>PASSED (EXCEEDED)</b> | Intent Accuracy <b>90.8%</b> , Fluency Recall <b>92.0%</b> , Urgency Safety <b>85.0%</b> using LLM few-shot prompts (Phase 1). |
| FR-02          | Route <b>security_incident</b> , <b>compliance_request</b> , <b>feature_request</b> , <b>unclear_request</b> safely. | <b>PASSED (EXCEEDED)</b> | <b>97.7% MNR Safety, 0 policy violations on security/compliance auto-send.</b>                                                 |

| Requirement ID | Requirement Promise                                                                          | Implementation Status | Empirical Evidence / Verification                                                                                                |
|----------------|----------------------------------------------------------------------------------------------|-----------------------|----------------------------------------------------------------------------------------------------------------------------------|
| FR-03          | Semantic search retrieval with 1-hop graph expansion from <code>documents</code> table only. | PASSED                | <b>81.2% answerable retrieval hit rate</b> (283/357 tickets), pre-attaching docs to 57.1% of escalations.                        |
| FR-04          | Grounded draft generation with explicit article citations.                                   | PASSED                | <b>0.0% Must-Not-Claim Rate, 100% placeholder string elimination, 91.9% pre-draft coverage.</b>                                  |
| FR-05          | Attach structured context packet to every escalated ticket.                                  | PASSED                | <b>100% of escalated tickets</b> carry structured context packets with pre-drafts and doc IDs.                                   |
| FR-06          | Present drafts to Tier 1/2 human agents for review before sending.                           | PASSED (EXCEEDED)     | <b>22.8% Copilot-Assisted FCR</b> (114/500 tickets), saving <b>~19 hours</b> of human drafting effort over 100% manual baseline. |
| FR-07          | Log 100% of automated routing decisions for auditability.                                    | PASSED                | <b>100% Decision Auditability</b> (500/500 decisions logged to SQLite <code>storage/decisions.db</code> ).                       |
| FR-08          | Normalise incoming tickets from all 4 channels (email, chat, docs_comment, forum).           | PASSED                | 100% parsing success across all 4 channels without text corruption.                                                              |
| FR-09          | Validate drafts for customer PII before release and block PII leakage.                       | PASSED                | <code>PiiValidator</code> active; <b>0 PII leaks</b> in outbound responses.                                                      |
| FR-10          | Produce performance metrics segmented by region, tier, and language fluency.                 | PASSED                | Metrics and fairness audit produced by <code>evaluation/harness.py</code> .                                                      |
| FR-11          | Restate core question from non-fluent English tickets in clear English.                      | PASSED                | <b>92.0% Non-Fluent Recall</b> , handling-time gap managed effectively.                                                          |
| FR-12          | Order agent review queue by SLA time remaining and intent urgency.                           | PASSED                | High-urgency response latency reduced from <b>478 min</b> to <b>21.6 sec.</b>                                                    |

| Requirement ID  | Requirement Promise                    | Implementation Status | Empirical Evidence / Verification                                                            |
|-----------------|----------------------------------------|-----------------------|----------------------------------------------------------------------------------------------|
| Target FCR      | Achieve FCR $\geq$ 65% benchmark.      | EXCEEDED (+18.8 pp)   | <b>83.8% Total Effective FCR</b> (61.0% Autonomous L1 AI + 22.8% Assisted L2 Copilot).       |
| Specialist Load | Reduce senior specialist queue burden. | EXCEEDED              | L3 Senior Specialist queue load collapsed from <b>281 to 81 tickets (-71.2% reduction)</b> . |

Related documents: [01 Discovery Workbook.pdf](#) · [02 Requirements PRD.pdf](#) · [03 Prompt Library.pdf](#) · [04 Sprint Plan.pdf](#) · [AyushYadav Effort Log.pdf](#)